

# Daniel Yarmoluk

AI Architect | Agentic AI and Enterprise Adoption | GitHub Copilot, MCP and A2A | Prompt Engineering and Governance

Minneapolis, MN | [daniel.yarmoluk@gmail.com](mailto:daniel.yarmoluk@gmail.com)

[linkedin.com/in/danyarmoluk](https://linkedin.com/in/danyarmoluk) | [github.com/Yarmoluk](https://github.com/Yarmoluk) | [pypi.org/user/danyarm](https://pypi.org/user/danyarm) | [huggingface.co/danyarm](https://huggingface.co/danyarm) | [graphifymd.com](https://graphifymd.com)

## PROFESSIONAL SUMMARY

---

AI architect who defines enterprise AI architecture, standards and governance, and builds the reference implementations teams work from. Fifteen years of Fortune 500 delivery including production-readiness, evaluation and Responsible AI governance frameworks for HIPAA-regulated environments. Hands-on across agentic AI: multi-step autonomous agent orchestration, prompt libraries, guardrails and adoption metrics. Built 6 production Model Context Protocol (MCP) services and 12 published Python packages — and MCP is GitHub Copilot's extension mechanism, configured per repository via `.vscode/mcp.json` and governed by the organisation-level MCP servers in Copilot policy, so the same servers extend Copilot agent mode unchanged. Deep RAG and vector search: built and benchmarked a full pipeline (FAISS and Chroma, sentence-transformers embeddings, chunking, top-k) against a typed knowledge graph across 44 domains and 7,758 queries. Multi-cloud across AWS, Microsoft Azure and Google Cloud. Patent pending. Adjunct professor of graduate AI systems.

## PROFESSIONAL EXPERIENCE

---

Graphify.md — Minneapolis, MN

*Founder and Principal Engineer, AI Infrastructure | 04/2026 – Present*

- Define enterprise AI architecture, standards, reference patterns and governance guidelines for adoption at scale — including Responsible AI practices, production-readiness gates and evaluation frameworks for regulated environments.
- Architect multi-step autonomous agent workflows and orchestration, and own the enterprise prompt strategy: a reusable prompt library of versioned, trigger-scoped templates (84 authored), zero-shot tool and output schemas, few-shot worked examples, and chain-of-thought decomposition made explicit through task-scoped subgraphs.
- Own adoption guardrails and productivity metrics end to end — authentication, license gating, per-method metering, rate limiting, and client-identity instrumentation. Rebuilt an analytics layer after proving its headline usage metric could not be attributed to any caller; naive adoption metrics measure plausibility, not outcomes.
- Provide hands-on Python reference implementations and integration patterns for teams — 12 published packages across 100+ releases, plus automated testing and CI/CD pipelines.
- Designed, built, deployed and operated 6 production Model Context Protocol (MCP) servers serving typed knowledge graphs to AI agents, and published 12 Python packages to PyPI across 100+ releases — owning the complete lifecycle from protocol design through incident response. All of it built solo, first package published within weeks of starting.
- Implemented the full MCP protocol surface in Python: JSON-RPC 2.0 initialize handshake with `clientInfo` and `protocolVersion` negotiation, `tools/list`, `tools/call`, `resources/list`, `prompts/list`, and notifications, with tool input and output schema validation.
- Diagnosed and fixed a production transport defect where every external MCP client received HTTP 421 Misdirected Request: DNS-rebinding protection remained active with an empty host allow-list behind a reverse proxy, rejecting all traffic before it reached a tool. Implemented explicit transport security settings with configurable allowed hosts.
- Built the metering and monetization layer for MCP: classified protocol methods as metered versus unmetered so the three-request handshake no longer consumed a client's entire free tier, implemented time-windowed free access with a call ceiling, per-IP and per-/24-subnet rate limiting, license key validation, and HTTP 402 Payment Required responses with signed payment challenges.
- Instrumented MCP servers with product analytics using PostHog's MCP event schema — capturing tool calls, client identity from the initialize handshake, latency, errors and funnel conversion — and corrected three instrumentation defects that had made usage unattributable to any caller.
- Implemented Agent-to-Agent (A2A) and machine-payment integrations: x402 / HTTP 402 request flow with EIP-3009 signed authorizations, Coinbase Developer Platform facilitator authentication, Base network settlement, and ERC-8004 agent identity — and documented the honest limitation that no mainnet facilitator existed for the required scheme.
- Built a conjunctive query operator for knowledge graphs exposed as a single MCP tool — returning concepts reachable from every anchor as an exact set intersection over declared edges, with typed relation paths, AND/OR modes, and a per-step cardinality trace that makes retrieval auditable.
- Designed a SHA-256 per-node source provenance system so any answer can be traced to the exact bytes of its source, including canonical-text hashing for CDN-rendered documents whose raw bytes carry rotating build tokens, and a verification tool that reports honestly when a source cannot be independently verified.
- Benchmarked Compressed Knowledge Graph retrieval against RAG and Microsoft GraphRAG across 44 domains and 7,758 queries: macro-F1 0.471 versus 0.123 (RAG) and 0.120 (GraphRAG), a 3.8x improvement; 269 versus 2,982 tokens per query, an 11x reduction; multi-hop F1 rising to 0.772 at 5 hops while RAG remained flat near 0.17. Published the dataset and harness openly.

- Built AI evaluation infrastructure: a reproducible retrieval-evaluation harness plus an independent evaluation against FrontierFinance, Samaya Research's open finance benchmark (220 queries, 11,543 rubric items, LLM judge panel) — evaluating against a third-party rubric rather than only a self-authored one.
- Authored and maintained 300+ domain knowledge graphs with typed relations and per-node source provenance across healthcare, financial regulation, data engineering, and AI infrastructure.
- Patent pending — three provisional applications filed 2026 covering knowledge graph architecture, encoding pipeline, and context infrastructure.

Slalom Consulting — Minneapolis, MN

*Principal — Solution Owner, AI and Digital Transformation | 09/2023 – 04/2026*

- Architected enterprise AI and digital transformation programs for Fortune 500 clients including UnitedHealthcare (clinical AI pipelines), Best Buy (retail AI), and Cargill (supply chain AI).
- Led cross-functional AI delivery across engineering, data science and business stakeholders, translating strategy into production system architecture and deployment roadmaps.
- Designed AI workflow orchestration and agent coordination patterns for enterprise deployment, including production readiness, monitoring and evaluation frameworks for regulated environments.
- Managed technical vendor relationships and platform integration across Anthropic, OpenAI, Azure OpenAI Service on Microsoft Azure AI, and AWS SageMaker.

West Monroe Partners — Minneapolis, MN

*Lead Solution Architect, Industrial IoT and Commercial AI | 2021 – 2023*

- Led AI and Internet of Things (IoT) architecture for industrial and commercial clients: production machine learning pipelines, real-time data systems, and multi-system orchestration.
- Designed production deployment architecture for AI systems in operational environments, covering reliability, monitoring and performance under load.

ATEK Access Technologies — Minneapolis, MN

*IoT and Data Science Practice Leader | 2016 – 2020*

- Built the data science and IoT practice from zero, including AWS, Siemens and IBM strategic partnerships and production machine learning deployment in industrial field operations.
- Led sensor data orchestration across distributed field devices, including real-time data pipeline design and system integration architecture.

University of St. Thomas — Minneapolis, MN

*Adjunct Professor, Graduate AI Systems (SEIS 666) | 2018 – Present*

- Eight years accelerating large language model (LLM) and knowledge graph capability for senior product and engineering leaders. Students are practicing technology executives and staff-level engineers who apply the material directly to live production initiatives at their own companies.
- Designed and continually rewrite graduate curriculum spanning production LLM systems, agent orchestration design, knowledge graph infrastructure, enterprise AI evaluation, and digital transformation — updated each term as the capability frontier moves.
- Enterprise AI adoption and enablement at leadership level: technical upskilling of senior technologists, change management, and closing the gap between what a system can do and what an organization will actually put into production.

## BENCHMARK — PUBLISHED AND REPRODUCIBLE

---

Compressed Knowledge Graph versus Retrieval-Augmented Generation (RAG) and Microsoft GraphRAG. Version 0.6.2, locked: 44 domains, 7,758 queries. Dataset and evaluation harness published openly (CC-BY-4.0) so results can be independently re-run.

- Answer quality: macro-F1 0.471 (CKG) versus 0.123 (RAG) and 0.120 (GraphRAG) — 3.8x.
- Context efficiency: 269 versus 2,982 tokens per query — 11x fewer tokens.
- Multi-hop reasoning: F1 rises to 0.772 at 5-hop depth while RAG stays flat near 0.170.
- Auditability: every answer traces to a declared relationship and, where source provenance is present, to a SHA-256 hash of the source bytes.
- Reproducibility: [github.com/Yarmoluk/ckg-benchmark](https://github.com/Yarmoluk/ckg-benchmark) and [huggingface.co/danyarm](https://huggingface.co/danyarm) — clone and re-run.

## PUBLICATIONS AND RECOGNITION

---

- Publication in preparation: a systematic benchmark of compressed knowledge graphs for enterprise AI, in collaboration with Dan McCreary, former Head of AI at TigerGraph.
- Media: featured in CIO Dive on enterprise context limitations; guest on the AI and Data Horizons podcast; host of All Things Data (1M+ downloads). Benchmark cited independently at Data Tech 2026.

## EDUCATION

---

- Graduate Degree, Data Science — University of St. Thomas
- Master of Business Administration — Loyola University Chicago
- Bachelor of Arts, International Affairs — University of Texas